

Information Retrieval

ARNAB BHATTACHARYA

arnabb@cse.iitk.ac.in

Department of Computer Science and Engineering,
Indian Institute of Technology, Kanpur,
India

14th June, 2017

ACM Summer School on NLP and ML

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
 - Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
- Inverse Document Frequency
- Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Definition

- What is **information retrieval**?

Definition

- What is **information retrieval**?

Information Retrieval

Retrieval (finding) of **information** (e.g., documents) that is mostly **unstructured** (e.g., text) and is **relevant** to a particular **need** (query) from a **large** collection

Definition

- What is **information retrieval**?

Information Retrieval

Retrieval (finding) of **information** (e.g., documents) that is mostly **unstructured** (e.g., text) and is **relevant** to a particular **need** (query) from a **large** collection

- Started with documents
- Has now extended to music, images, graphs

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
 - Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i,j}\}$

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i_j}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i,j}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**
- *Information retrieval task*
 - Given a set of terms from the vocabulary, find the documents that contain them

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i,j}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**
- *Information retrieval task*
 - Given a set of terms from the vocabulary, find the documents that contain them
 - Sense is generally AND; may sometimes involve OR and NOT

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i,j}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**
- *Information retrieval task*
 - Given a set of terms from the vocabulary, find the documents that contain them
 - Sense is generally AND; may sometimes involve OR and NOT
- Unix utility `grep` solves it nicely

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{ij}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**
- *Information retrieval task*
 - Given a set of terms from the vocabulary, find the documents that contain them
 - Sense is generally AND; may sometimes involve OR and NOT
- Unix utility `grep` solves it nicely
- Boolean bit matrix for each document and each term
- Boolean algebra on the bit vectors of the queried terms

Document Retrieval Task

- Assume a set $\mathcal{D} = \{D_1, \dots, D_m\}$ of *documents*
- Each D_i is composed of a number of *terms* $D_i = \{t_{i,j}\}$
- The total collection of terms is called the **dictionary** or **lexicon** or **vocabulary**
- *Information retrieval task*
 - Given a set of terms from the vocabulary, find the documents that contain them
 - Sense is generally AND; may sometimes involve OR and NOT
- Unix utility `grep` solves it nicely
- Boolean bit matrix for each document and each term
- Boolean algebra on the bit vectors of the queried terms
- *Not* scalable (remember “large” collections)

Inverted Index


- “Invert” the sense

Inverted Index

- “Invert” the sense – let terms be composed of documents


Inverted Index

- “Invert” the sense – let terms be composed of documents
- Each term contains a list of documents (called **postings**) that it appears in
 - May include the number of times it appears in a document, called **term frequency**
- This list of documents is called a **postings list**
 - Its size, called **document frequency**, is the number of documents a term appears
- The set of lists for all the terms is called **postings**
 - Its size is size of the vocabulary



Inverted Index

- “Invert” the sense – let terms be composed of documents
- Each term contains a list of documents (called **postings**) that it appears in
 - May include the number of times it appears in a document, called **term frequency**
- This list of documents is called a **postings list**
 - Its size, called **document frequency**, is the number of documents a term appears
- The set of lists for all the terms is called **postings**
 - Its size is size of the vocabulary



- Postings list is maintained as a linked list

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$
- If lists are sorted by document ids
 - Merging as in mergesort
 - Two pointers to walk along the lists
 - Time complexity is $O(x + y)$

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$
- If lists are sorted by document ids
 - Merging as in mergesort
 - Two pointers to walk along the lists
 - Time complexity is $O(x + y)$
- For multiple query terms: “ancient”, “Indian” and “system”

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$
- If lists are sorted by document ids
 - Merging as in mergesort
 - Two pointers to walk along the lists
 - Time complexity is $O(x + y)$
- For multiple query terms: “ancient”, “Indian” and “system”
 - Start with terms having *smallest* postings lists
 - Progressively merge
 - Size of final answer set is never more than any intermediate step

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$
- If lists are sorted by document ids
 - Merging as in mergesort
 - Two pointers to walk along the lists
 - Time complexity is $O(x + y)$
- For multiple query terms: “ancient”, “Indian” and “system”
 - Start with terms having *smallest* postings lists
 - Progressively merge
 - Size of final answer set is never more than any intermediate step
- OR is more time consuming

Querying using Inverted Index

- Find documents where terms “Indian” and “system” occur
- Get corresponding postings lists
- *Intersect* them
- For two lists of size x and y , time required is $O(x.y)$
- If lists are sorted by document ids
 - Merging as in mergesort
 - Two pointers to walk along the lists
 - Time complexity is $O(x + y)$
- For multiple query terms: “ancient”, “Indian” and “system”
 - Start with terms having *smallest* postings lists
 - Progressively merge
 - Size of final answer set is never more than any intermediate step
- OR is more time consuming
- Query with only NOT is impractical for large collections
 - All right with AND: “system” AND NOT “Indian”

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 **Tokenization**
 - **Stopwords**
 - **Token Normalization**
- 4 Scoring
 - Zone
 - Term Frequency
- Inverse Document Frequency
- Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Tokenization

- **Terms** are useful semantic units that need to be indexed for searching
- **Tokenization** is the process of breaking the text into *terms*

Tokenization

- **Terms** are useful semantic units that need to be indexed for searching
- **Tokenization** is the process of breaking the text into *terms*
 - Sounds obvious in English with whitespaces
 - What about Indian languages with *sandhi* and *samash*?
 - What about CJK languages with no space between words (or sometimes even paragraphs)?
 - Non-trivial for even some European languages
 - “computerlinguistik” (German, computational linguistics)
 - “l’ensemble” (French, the collection)
- **Word segmentation**

Tokenization

- **Terms** are useful semantic units that need to be indexed for searching
- **Tokenization** is the process of breaking the text into *terms*
 - Sounds obvious in English with whitespaces
 - What about Indian languages with *sandhi* and *samash*?
 - What about CJK languages with no space between words (or sometimes even paragraphs)?
 - Non-trivial for even some European languages
 - “computerlinguistik” (German, computational linguistics)
 - “l’ensemble” (French, the collection)
- **Word segmentation**

मूर्तमहेश्वरमुज्ज्वलभास्करमिष्टममरनरवन्दम्

Tokenization

- **Terms** are useful semantic units that need to be indexed for searching
- **Tokenization** is the process of breaking the text into *terms*
 - Sounds obvious in English with whitespaces
 - What about Indian languages with *sandhi* and *samash*?
 - What about CJK languages with no space between words (or sometimes even paragraphs)?
 - Non-trivial for even some European languages
 - “computerlinguistik” (German, computational linguistics)
 - “l’ensemble” (French, the collection)
- **Word segmentation**

मूर्तमहेश्वरमुज्ज्वलभास्करमिष्टममरनरवन्दम्

मूर्त + महेश्वरम् + उज्ज्वल + भास्करम् + इष्टम् + अमर + नर + वन्दम्

Problems in English

- Idioms

Problems in English

- Idioms
 - “hot potato”, “couch potato”

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation
 - “co-education”

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation
 - “co-education”
- White space

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation
 - “co-education”
- White space
 - “New Delhi”

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation
 - “co-education”
- White space
 - “New Delhi”
- Combinations

Problems in English

- Idioms
 - “hot potato”, “couch potato”
- Word markers or punctuations
 - “can’t”, “O’Neille”
- Hyphenation
 - “co-education”
- White space
 - “New Delhi”
- Combinations
 - “isn’t New Delhi-Uttar Pradesh a good example?”

Stopwords

- Some terms add little *semantic* content to search
 - “to”, “and”, “the”
- These are called **stopwords**

Stopwords

- Some terms add little *semantic* content to search
 - “to”, “and”, “the”
- These are called **stopwords**
- Terms that have a very high *frequency* are candidates
- Generally, manually polished by linguistics

Stopwords

- Some terms add little *semantic* content to search
 - “to”, “and”, “the”
- These are called **stopwords**
- Terms that have a very high *frequency* are candidates
- Generally, manually polished by linguistics
- *Variety* in the **corpus** of documents is very important

Stopwords

- Some terms add little *semantic* content to search
 - “to”, “and”, “the”
- These are called **stopwords**
- Terms that have a very high *frequency* are candidates
- Generally, manually polished by linguistics
- *Variety* in the **corpus** of documents is very important
- If only newspaper articles,
 - “police” appears with very high frequency

Stopwords

- Some terms add little *semantic* content to search
 - “to”, “and”, “the”
- These are called **stopwords**
- Terms that have a very high *frequency* are candidates
- Generally, manually polished by linguistics
- *Variety* in the **corpus** of documents is very important
- If only newspaper articles,
 - “police” appears with very high frequency
- Removing stopwords from queries may sometimes be erroneous
 - “to be or not to be”
- Therefore, web search engines do not bother to remove stopwords

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**
- **Stemming** or **lemmatization** refers to stripping the word to its **root** or **lemma**
 - “system”, “systems”, “systematic”
 - Requires **morphological analysis** and is language specific

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**
- **Stemming** or **lemmatization** refers to stripping the word to its **root** or **lemma**
 - “system”, “systems”, “systematic”
 - Requires **morphological analysis** and is language specific
- Synonyms
 - “query”, “question”

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**
- **Stemming** or **lemmatization** refers to stripping the word to its **root** or **lemma**
 - “system”, “systems”, “systematic”
 - Requires **morphological analysis** and is language specific
- Synonyms
 - “query”, “question”
- Spelling versions
 - “color”, “colour”

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**
- **Stemming** or **lemmatization** refers to stripping the word to its **root** or **lemma**
 - “system”, “systems”, “systematic”
 - Requires **morphological analysis** and is language specific
- Synonyms
 - “query”, “question”
- Spelling versions
 - “color”, “colour”
- Token normalization finds more documents
 - Increases recall but decreases precision

Token Normalization

- Should documents containing “systems” be retrieved for “system”?
- Tokens are divided into **equivalence classes**
- **Stemming** or **lemmatization** refers to stripping the word to its **root** or **lemma**
 - “system”, “systems”, “systematic”
 - Requires **morphological analysis** and is language specific
- Synonyms
 - “query”, “question”
- Spelling versions
 - “color”, “colour”
- Token normalization finds more documents
 - Increases recall but decreases precision
- Character set is important
 - “pena” (sorrow) and “peña” (cliff) in Spanish

Documents

- What are documents?

Documents

- What are **documents**?
 - Entire book may or may not be useful
 - May be chapters or sections or sentences

Documents

- What are **documents**?
 - Entire book may or may not be useful
 - May be chapters or sections or sentences
- Depends on how **context** is defined

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
 - Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Non-Boolean Match

- A query may be more general than just a set of terms
- It may itself be another document or some free text
- No document can be then expected to match it fully

Non-Boolean Match

- A query may be more general than just a set of terms
- It may itself be another document or some free text
- No document can be then expected to match it fully
- *Information retrieval task*
 - Given a free text query, find the most **similar** documents

Non-Boolean Match

- A query may be more general than just a set of terms
- It may itself be another document or some free text
- No document can be then expected to match it fully
- *Information retrieval task*
 - Given a free text query, find the most **similar** documents
- “Similarity” of a document d with the query q can be measured by a **score** $s(d, q)$

Non-Boolean Match

- A query may be more general than just a set of terms
- It may itself be another document or some free text
- No document can be then expected to match it fully
- *Information retrieval task*
 - Given a free text query, find the most **similar** documents
- “Similarity” of a document d with the query q can be measured by a **score** $s(d, q)$
- *Information retrieval task*
 - Given a free text query, find the documents that have the largest scores

Zones




- Each document is generally associated with **metadata**, e.g., title, author, date, etc.
- Sometimes, queries on indexes of these fields, called **parametric indexes**, are also asked
 - Find documents where “Indian” appears in the title

Zones

- Each document is generally associated with **metadata**, e.g., title, author, date, etc.
- Sometimes, queries on indexes of these fields, called **parametric indexes**, are also asked
 - Find documents where “Indian” appears in the title
- Fields are generalized to **zones** that may contain free text as well

Zones

- Each document is generally associated with **metadata**, e.g., title, author, date, etc.
- Sometimes, queries on indexes of these fields, called **parametric indexes**, are also asked
 - Find documents where “Indian” appears in the title
- Fields are generalized to **zones** that may contain free text as well
- Separate inverted indexes can be built for each zone
- Or, zone may be mentioned *explicitly* in a single inverted index



Weighted Zone Scoring

- Each zone has a weight that adds up to 1
 - “title” has 0.3, and “body” has 0.7

Weighted Zone Scoring

- Each zone has a weight that adds up to 1
 - “title” has 0.3, and “body” has 0.7
- Given a query term, each zone scores a **zone score**
- It is 1 if the term is inside the zone; 0 otherwise

Weighted Zone Scoring

- Each zone has a weight that adds up to 1
 - “title” has 0.3, and “body” has 0.7
- Given a query term, each zone scores a **zone score**
- It is 1 if the term is inside the zone; 0 otherwise
- **Weighted zone score** is the linear sum of the zone scores

$$s(q, d) = \sum_{\forall z \in d} s_z(q)$$

Weighted Zone Scoring

- Each zone has a weight that adds up to 1
 - “title” has 0.3, and “body” has 0.7
- Given a query term, each zone scores a **zone score**
- It is 1 if the term is inside the zone; 0 otherwise
- **Weighted zone score** is the linear sum of the zone scores

$$s(q, d) = \sum_{\forall z \in d} s_z(q)$$

- Given a Boolean combination of query terms, each zone is scored
- These scores are then accumulated

Weighted Zone Scoring

- Each zone has a weight that adds up to 1
 - “title” has 0.3, and “body” has 0.7
- Given a query term, each zone scores a **zone score**
- It is 1 if the term is inside the zone; 0 otherwise
- **Weighted zone score** is the linear sum of the zone scores

$$s(q, d) = \sum_{\forall z \in d} s_z(q)$$

- Given a Boolean combination of query terms, each zone is scored
- These scores are then accumulated
- Weights of zones
 - Can be supplied by the application
 - Machine learned

Term Frequency

- Moving away from the binary model
- If a document contains a query term more number of times, it is more important and should score higher
- Weight of a document d is, therefore, simply the number of times the term t appears in it, called the **term frequency**

$$tf(t, d) = |t \in d|$$

Term Frequency

- Moving away from the binary model
- If a document contains a query term more number of times, it is more important and should score higher
- Weight of a document d is, therefore, simply the number of times the term t appears in it, called the **term frequency**

$$tf(t, d) = |t \in d|$$

- This assumes the **bag of words** model
- *Context* and *sequence* are lost
 - I love butter but I hate cheese
 - I love cheese but I hate butter

Inverse Document Frequency

- Certain terms may appear across all or most documents
 - “bat” in cricket pages
- Consequently, they discriminate little among the documents and are not useful

Inverse Document Frequency

- Certain terms may appear across all or most documents
 - “bat” in cricket pages
- Consequently, they discriminate little among the documents and are not useful
- **Document frequency** of a term is the number of documents it appears in
 - Lesser is more discriminative

Inverse Document Frequency

- Certain terms may appear across all or most documents
 - “bat” in cricket pages
- Consequently, they discriminate little among the documents and are not useful
- **Document frequency** of a term is the number of documents it appears in
 - Lesser is more discriminative
- If m is the total number of documents in the corpus

$$idf(t) = \log \frac{m}{df_t}$$

- This is called the **inverse document frequency**

Inverse Document Frequency

- Certain terms may appear across all or most documents
 - “bat” in cricket pages
- Consequently, they discriminate little among the documents and are not useful
- **Document frequency** of a term is the number of documents it appears in
 - Lesser is more discriminative
- If m is the total number of documents in the corpus

$$idf(t) = \log \frac{m}{df_t}$$

- This is called the **inverse document frequency**
- Logarithmic to make it less drastic

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents
 - *Low* when t appears in many documents

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents
 - *Low* when t appears in many documents
 - *High* when t appears many number of times in d

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents
 - *Low* when t appears in many documents
 - *High* when t appears many number of times in d
 - *Low* when t appears few number of times in d

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents
 - *Low* when t appears in many documents
 - *High* when t appears many number of times in d
 - *Low* when t appears few number of times in d
 - *Zero* if t does not appear at all in d

Tf-idf

- Combination of term frequency (**tf**) and inverse document frequency (**idf**)
- Weight of a term t in a document d is

$$tf-idf(t, d) = tf(t, d) \times idf(t)$$

- The **tf-idf** score has following properties
 - *High* when t appears in a small number of documents
 - *Low* when t appears in many documents
 - *High* when t appears many number of times in d
 - *Low* when t appears few number of times in d
 - *Zero* if t does not appear at all in d
- Tf-idf has many different forms

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
 - Inverse Document Frequency
 - Tf-idf
- 5 Document Vector**
 - **Vector Model**
 - **Document Similarity**
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Document Vector

- Each document d has a score with each term t in the vocabulary
 - If t is absent in d , then this score is 0

Document Vector

- Each document d has a score with each term t in the vocabulary
 - If t is absent in d , then this score is 0
- Imagine a n -dimensional vector space where n is the total number of terms in the vocabulary
- Each document can be, thus, thought of as a vector (point) in this n -dimensional space
- Its coordinates are the scores corresponding to the scores

$$d[t_i] = tf-idf(t_i, d)$$

- This is called the **document vector** model

Exercise

- d_1 : Water, water everywhere, not a drop to drink
- d_2 : I have filtered water
- d_3 : Drinking and driving is not good
- d_4 : Water quality is not good here
- d_5 : Milk is not good for health
- d_6 : Drinking water just after dinner is not healthy

Exercise

- d_1 : Water, water everywhere, not a drop to drink
- d_2 : I have filtered water
- d_3 : Drinking and driving is not good
- d_4 : Water quality is not good here
- d_5 : Milk is not good for health
- d_6 : Drinking water just after dinner is not healthy
- Query q : drinkable water

Exercise

- d_1 : Water, water everywhere, not a drop to drink
- d_2 : I have filtered water
- d_3 : Drinking and driving is not good
- d_4 : Water quality is not good here
- d_5 : Milk is not good for health
- d_6 : Drinking water just after dinner is not healthy
- Query q : drinkable water
- Find tf, idf (with \log_2) and tf-idf (\log_2) scores

Similarity between Documents

- What is the “similarity” between two documents (i.e., their vectors)?

Similarity between Documents

- What is the “similarity” between two documents (i.e., their vectors)?
- *Euclidean distance* may not be suitable
 - Longer documents have larger distances

Cosine Similarity

- Consider two documents d_1 and d_2 with their corresponding document vectors $\vec{V}(d_1)$ and $\vec{V}(d_2)$
- **Cosine similarity** measures the *normalised dot product*

$$\text{sim}(d_1, d_2) = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| \cdot |\vec{V}(d_2)|}$$

- Measures the cosine of the angle between the vectors

Cosine Similarity

- Consider two documents d_1 and d_2 with their corresponding document vectors $\vec{V}(d_1)$ and $\vec{V}(d_2)$
- Cosine similarity** measures the *normalised dot product*

$$\text{sim}(d_1, d_2) = \frac{\vec{V}(d_1) \cdot \vec{V}(d_2)}{|\vec{V}(d_1)| \cdot |\vec{V}(d_2)|}$$

- Measures the cosine of the angle between the vectors
- Consider the **length-normalised** document vectors

$$\vec{v}(d_i) = \frac{\vec{V}(d_i)}{|\vec{V}(d_i)|}$$

- Then, cosine similarity is their dot product

$$\text{sim}(d_1, d_2) = \vec{v}(d_1) \cdot \vec{v}(d_2)$$

Example

Term	d_1	d_2	d_3
Indian	115	58	20
ancient	10	7	11
system	2	0	6

Example

Term	d_1	d_2	d_3
Indian	115	58	20
ancient	10	7	11
system	2	0	6
length	115.45	58.42	23.60

Example

Term	d_1	d_2	d_3
Indian	115	58	20
ancient	10	7	11
system	2	0	6
length	115.45	58.42	23.60

- Similarities between documents

$$\text{sim}(d_1, d_2) = \frac{115}{115.45} \cdot \frac{58}{58.42} + \frac{10}{115.45} \cdot \frac{7}{58.42} + \frac{2}{115.45} \cdot \frac{0}{58.42} = 0.99$$

- d_1 and d_2 is the closest pair

Query Vector

- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents

Query Vector

- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents
- Consider query q with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$ with $\vec{v}(q) = (0.00, 0.71, 0.71)$

Query Vector

- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents
- Consider query q with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$ with $\vec{v}(q) = (0.00, 0.71, 0.71)$
- Then, most similar document is the one with the *highest* cosine similarity

Query Vector

- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents
- Consider query q with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$ with $\vec{v}(q) = (0.00, 0.71, 0.71)$
- Then, most similar document is the one with the *highest* cosine similarity
- In previous example, $\text{sim}(q, d_3)$ is highest
- Thus, d_3 is the most similar document

Query Vector

- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents
- Consider query q with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$ with $\vec{v}(q) = (0.00, 0.71, 0.71)$
- Then, most similar document is the one with the *highest* cosine similarity
- In previous example, $\text{sim}(q, d_3)$ is highest
- Thus, d_3 is the most similar document
- If $\vec{v}(q) \cdot \vec{v}(d_i)$ is the highest, then $\vec{V}(q) \cdot \vec{v}(d_i)$ is the highest as well
 - Normalization of query is not required

Query Vector


- Similar to a document, the query can also be viewed as a vector
- It is again just a *bag of words*
- Find similarities with documents and retrieve the top- k documents
- Consider query q with the keywords “ancient” and “system”
- $\vec{V}(q) = (0, 1, 1)$ with $\vec{v}(q) = (0.00, 0.71, 0.71)$
- Then, most similar document is the one with the *highest* cosine similarity
- In previous example, $\text{sim}(q, d_3)$ is highest
- Thus, d_3 is the most similar document
- If $\vec{v}(q) \cdot \vec{v}(d_i)$ is the highest, then $\vec{V}(q) \cdot \vec{v}(d_i)$ is the highest as well
 - Normalization of query is not required
- Brute-force method of computing scores with all the documents and ranking them is *not* scalable

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
- 5 Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 Conclusions

Skip Lists


- **Skip lists** are used to traverse linked lists faster



- “Skips” are provided as jumps

Skip Lists


- **Skip lists** are used to traverse linked lists faster



- “Skips” are provided as jumps
- Useful when taking intersection of postings lists
- Once 41 is reached in the second list, 16 can jump to 28 in the first list

Skip Lists


- **Skip lists** are used to traverse linked lists faster



- “Skips” are provided as jumps
- Useful when taking intersection of postings lists
- Once 41 is reached in the second list, 16 can jump to 28 in the first list
- How to determine skip positions?

Skip Lists


- **Skip lists** are used to traverse linked lists faster



- “Skips” are provided as jumps
- Useful when taking intersection of postings lists
- Once 41 is reached in the second list, 16 can jump to 28 in the first list
- How to determine skip positions?
- Look at closely occurring values; insert a skip to the end

Skip Lists

- **Skip lists** are used to traverse linked lists faster



- “Skips” are provided as jumps
- Useful when taking intersection of postings lists
- Once 41 is reached in the second list, 16 can jump to 28 in the first list
- How to determine skip positions?
- Look at closely occurring values; insert a skip to the end
- \sqrt{l} equally spaced skips for a l -length list

Approximation

- Finding the *exact* top- k documents is too difficult
- Can resort to *approximate* top- k most relevant documents

Approximation

- Finding the *exact* top- k documents is too difficult
- Can resort to *approximate* top- k most relevant documents
- Notion of similarity is not perfect anyway
 - Cosine similarity is simply one measure

Approximation

- Finding the *exact* top- k documents is too difficult
- Can resort to *approximate* top- k most relevant documents
- Notion of similarity is not perfect anyway
 - Cosine similarity is simply one measure
- Information needs are generally satisfied with approximate answers
 - Google search

Approximation

- Finding the *exact* top- k documents is too difficult
- Can resort to *approximate* top- k most relevant documents
- Notion of similarity is not perfect anyway
 - Cosine similarity is simply one measure
- Information needs are generally satisfied with approximate answers
 - Google search
- Retrieval process becomes much faster

Approximation

- Finding the *exact* top- k documents is too difficult
- Can resort to *approximate* top- k most relevant documents
- Notion of similarity is not perfect anyway
 - Cosine similarity is simply one measure
- Information needs are generally satisfied with approximate answers
 - Google search
- Retrieval process becomes much faster
- Generally, a two-step process
 - 1 Retrieve *approximate* top- K documents where $k \leq K \ll m$
 - 2 Retrieve *exact* top- k from K

Improving Efficiency

Improving Efficiency

- Only terms that appear in query need to be examined
 - Rest of the scores are 0

Improving Efficiency

- Only terms that appear in query need to be examined
 - Rest of the scores are 0
- Filter query terms whose *idf* is too low
 - Similar to stop-word idea
 - “to be or not to be”

Improving Efficiency

- Only terms that appear in query need to be examined
 - Rest of the scores are 0
- Filter query terms whose *idf* is too low
 - Similar to stop-word idea
 - “to be or not to be”
- *Pre-compute* **champion lists** for each term
 - Documents ranked for only that term
 - Offline process
 - Take union of top- r of every query term to get top- K

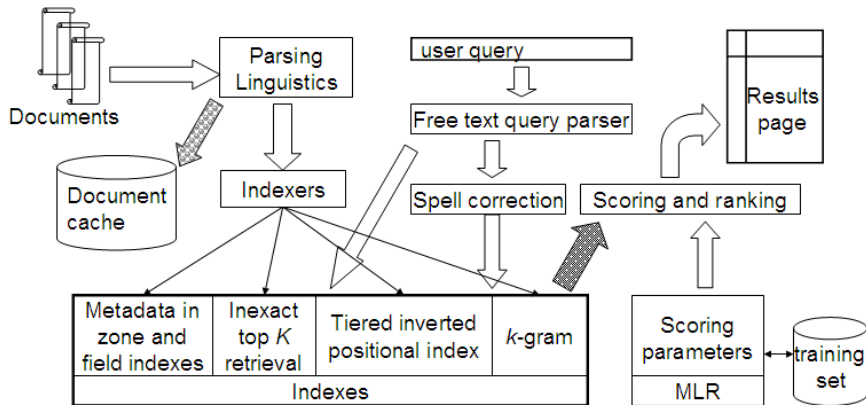
Improving Efficiency

- Only terms that appear in query need to be examined
 - Rest of the scores are 0
- Filter query terms whose *idf* is too low
 - Similar to stop-word idea
 - “to be or not to be”
- *Pre-compute champion lists* for each term
 - Documents ranked for only that term
 - Offline process
 - Take union of top- r of every query term to get top- K
- Build *tiered index*
 - Each level (tier) lists only those documents whose *tf* for the term is greater than a threshold
 - Continue with tiers till top- K results are obtained

Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
- 5 Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 **System**
- 8 Conclusions

The Complete Information Retrieval System



Outline

- 1 Introduction
 - Motivation
- 2 Document Retrieval
 - Inverted Index
- 3 Tokenization
 - Stopwords
 - Token Normalization
- 4 Scoring
 - Zone
 - Term Frequency
 - Inverse Document Frequency
 - Tf-idf
- 5 Document Vector
 - Vector Model
 - Document Similarity
- 6 Scalability
 - Skip List
 - Inexact Retrieval
- 7 System
- 8 **Conclusions**

Conclusions

- Inverted Index
- Tokenization
- Term Frequency, Document Frequency (Tf-idf)
- Document Vector
- Document Scoring

Conclusions

- Inverted Index
- Tokenization
- Term Frequency, Document Frequency (Tf-idf)
- Document Vector
- Document Scoring

THANK YOU!

Conclusions

- Inverted Index
- Tokenization
- Term Frequency, Document Frequency (Tf-idf)
- Document Vector
- Document Scoring

THANK YOU!

Questions?
Answers!